Causal Reliable Broadcast

Seif Haridi – Royal Institute of Technology
Peter Van Roy – Université catholique de Louvain

haridi(at)kth.se
peter.vanroy(at)uclouvain.be

Motivation

- Assume we have a chat application
  - Whatever written is reliably broadcast to group
- If you get the following output, is it ok?

[ali]	Are you sure? Not E?
[MrsY]	Auditorium D at Forum
[MrX]	Does anyone know where the lecture is today?

- MrX's message caused MrsY's message,
  - MrsY's message caused Ali's message

Motivation (2)

- Does uniform reliable broadcast remedy this? [d]
- Causal reliable broadcast solves this
  - Deliveries in causal order!
- Causality is same as happened-before relation by Lamport!

Causality Recalled

- Let m_1 and m_2 be any two messages: $m_1 \rightarrow m_2$ (m_1 causally precedes m_2) if
  - ***C1 (FIFO order)***.
    - Some process p_i broadcasts m_1 before broadcasting m_2
  - ***C2 (Network order)***.
    - Some process p_i delivers m_1 and later broadcasts m_2
  - ***C3 (Transitivity)***.
    - There is a message m' such that $m_1 \rightarrow m'$ and $m' \rightarrow m_2$

Causal Broadcast Interface

- **Module:**
  - Name: CausalOrder (co)
- ***Events***
  - Request: ⟨coBroadcast | m⟩
  - Indication: ⟨coDeliver | src, m⟩
- ***Property:***
  - ***CB***: If node p_i delivers m_1, then p_i must have delivered every message causally preceding (→) m_1 before m_1
- Is this useful? How can it be satisfied? [d]
  - It is only safety. Satisfy it by never delivering!

Causality

- ***C1 (FIFO order)***.
  - Some process p_i broadcasts m_1 before broadcasting m_2

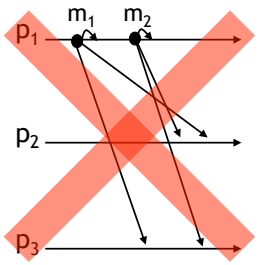

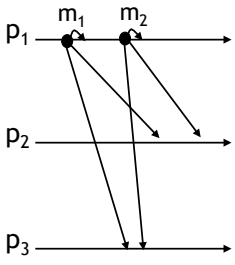

Causality (2)

- ***C2 (Network order)***.
  - Some process p_i delivers m_1 and later broadcasts m_2

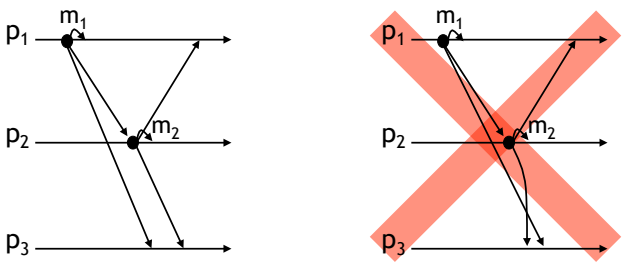

Causality (3)

- ***C3 (Transitivity)***.
  - There is a message m' such that $m_1 \rightarrow m'$ and $m' \rightarrow m_2$

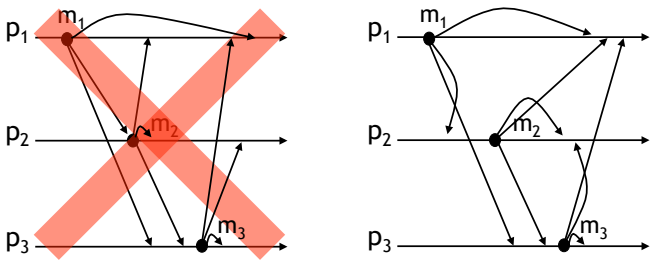

Different Causalities

- ***Property:***
  - ***CB***: If node p_i delivers m_1, then p_i must deliver every message causally preceding ($\rightarrow$) m_1 before m_1
  - ***CB'***: If p_j delivers m_1 and m_2, and $m_1 \rightarrow m_2$, then p_j must deliver m_1 before m_2
- What is the difference? [d]

Violates CB and CB'

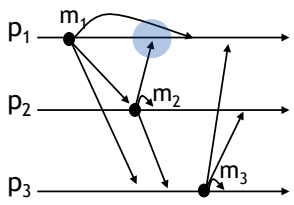

Violates CB, not CB'

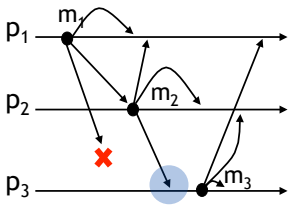

- Indeed, CB implies CB'

Reliable Causal Broadcast Interface

- **Module:**
  - Name: ReliableCausalOrder (rco)
- ***Events***
  - Request: ⟨rcoBroadcast | m⟩
  - Indication: ⟨rcoDeliver | src, m⟩
- ***Property:***
  - ***RB1-RB4*** from regular reliable broadcast
  - ***CB***: If node p_i delivers m, then p_i must deliver every message causally preceding ($\rightarrow$) m before m

Uniform Reliable Causal Broadcast

- **Module:**
  - Name: UniformReliableCausalOrder (urco)
- ***Events***
  - Request: ⟨urcoBroadcast | m⟩
  - Indication: ⟨urcoDeliver | src, m⟩
- ***Property:***
  - ***URB1-URB4*** from uniform reliable broadcast
  - ***CB***: If node p_i delivers m, then p_i must deliver every message causally preceding (→) m before m

Seif Haridi, based on Ali Ghodsi's slides

Idea Reuse…

- Reuse RB for CB
  - Use **reliable broadcast** abstraction to implement **reliable causal broadcast**
  - Use **uniform reliable broadcast** abstraction to implement **uniform causal broadcast**
- This gives a layered architecture!
  - CB component on top of RB component

Seif Haridi, based on Ali Ghodsi's slides

Towards an Implementation

- Main idea
  - Each broadcasted message carries a history
  - Before delivery, ensure causality
- First algorithm
  - History is set of all causally preceding messages

Fail-Silent No-Waiting Causal Bcast

- Each message m carries ordered list of causally preceding messages in $\mathbf{past_m}$
- Whenever a node rbDelivers m
  - coDeliver causally preceding messages in $\mathbf{past_m}$
  - coDelivers m
    - Avoid duplicates using **delivered** set

Execution (direct override)

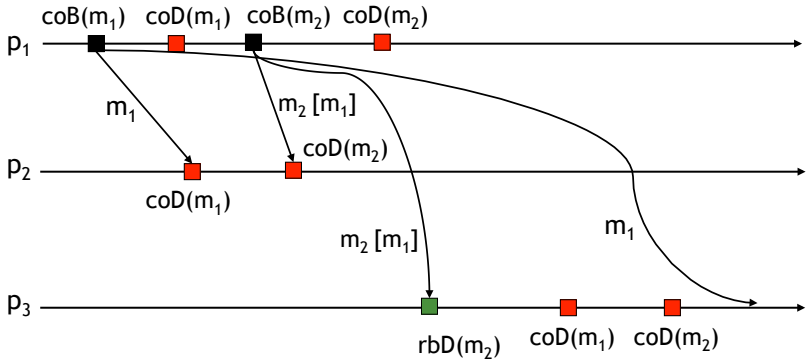

Execution (indirect override)

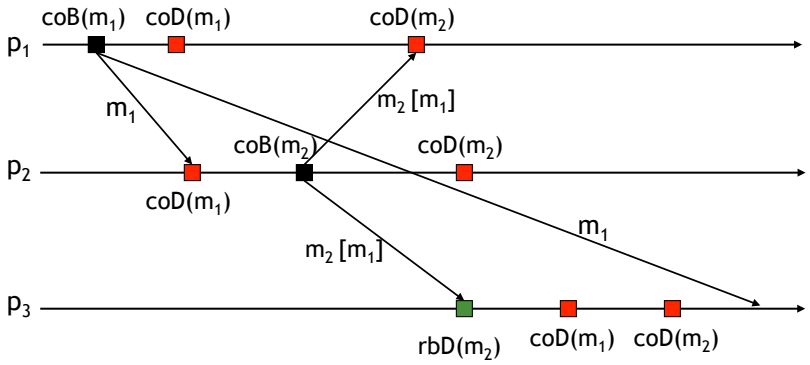

Fail-Silent Causal Broadcast Impl

- **Implements:**
  - ReliableCausalOrderBroadcast (rco).
- **Uses:** ReliableBroadcast (rb).
- **upon event** ⟨Init⟩ **do**
  - delivered := ∅; past := nil
- **upon event** ⟨rcoBroadcast | m⟩ **do**
  - **trigger** ⟨rbBroadcast | (DATA, past, m)⟩
  - past := append(past, <p_i, m>)

Append this message to past history

Fail-Silent Causal Broadcast Impl (2)

- **upon event** ⟨rbDeliver | pi,(DATA, $past_m$, m)⟩ **do**
  - **if** m∉delivered **then**
  - **forall** (s_n,n)∈$past_m$ **do** ← in ascending order
  - **if** n∉delivered **then**
  - **trigger** ⟨rcoDeliver|s_n, n⟩ ← deliver preceding messages
  - delivered := delivered∪{n}
  - past := append(past, <s_n,n>) ← append to history
  - **trigger** ⟨rcoDeliver|p_i,m⟩ ← deliver current message
  - delivered := delivered∪{m}
  - past := append(past, <p_i,m>) ← append to history

Correctness

- RB1-RB4 follow from use of RB
  - No creation and no duplication still satisfied
  - Validity still satisfied
    - Some messages might be delivered earlier, never later
  - Agreement directly from RB
- CO by induction on prefixes of executions
  - It is vacuously true for empty executions
  - Assume it is true for all deliveries of a prefix
    - Then it is true for any extension with one event

Improving the Algorithm

- Disadvantage of algorithm is that the message size (bit complexity) grows
- Useful idea
  - Garbage collect old messages
- Implementation of GC
  - Ack receipt of every message m to all
  - Use perfect failure detector P
    - Determine with P when all correct nodes got message m
    - Delete m from past when all correct nodes got m

GC Implementation

- **Uses:** ReliableBroadcast (rb), PerfectFailureDetector (P)
- **upon event** ⟨Init⟩ **do**
  - delivered := ∅; past := nil
  - correct := Π
  - forall m: ack[m] := ∅ ← **bookkeeping of acks**
- **upon event** ⟨crash | p_i⟩ **do**
  - correct := correct \ {p_i}
- **upon event** m∈delivered **and** self∉ack[m] **do** ← **called upon coDeliver**
  - ack := ack[m] U {self}
  - **trigger** ⟨rbBroadcast | (ACK, m)⟩ ← **ack to all**
- **upon event** ⟨rbDeliver | p_i, [ACK, m]⟩ **do**
  - ack := ack[m] U {p_i}
  - **if** correct⊆ack[m] **do**
  - past := remove(past, <x, m>) ← **When received ack from all, GC m from any x**

Seif Haridi, based on Ali Ghodsi's slides

GC Questions

- What about the acks? [d]
  - The ack[m] array also grows with time
  - How do we garbage collect it?
- What happens if we use ◊P? [d]
  - Does the garbage collector still work?

Seif Haridi, based on Ali Ghodsi's slides

Towards Another Implementation

- Main idea
  - Each broadcasted message carries a history
  - Before delivery, ensure causality
- First algorithm
  - History is set of all causally preceding messages
- Second algorithm [d]
  - History is a vector timestamp

Fail-Silent Waiting Causal Broadcast

- Represent past history by vector clock (VC)
- Slightly modify the VC implementation
  - At node p_i
    - VC[i]: number of messages p_i coBroadcasted
    - VC[j], j≠i: number of messages p_i coDelivered from p_j
- Idea: vector clock only for relevant events

Fail-Silent Waiting Causal Broadcast

- Upon CO broadcast m
  - Piggyback VC and RB broadcast m
- Upon RB delivery of m with attached VC_m compare VC_m with local VC_i
  - Only deliver m once VC_m precedes ($\leq$) VC_i

Execution

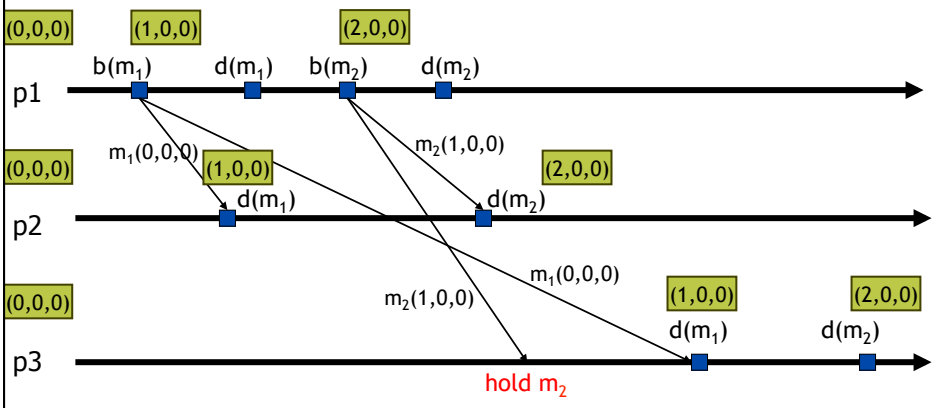

Fail-Silent Waiting Causal Impl.

- **Implements:**
  - ReliableCausalOrderBroadcast (rco)
- **Uses:** ReliableBroadcast (rb)

- **upon event** ⟨Init⟩ **do**
  - **forall** pi ∈ Π **do** VC[i] := 0

- **upon event** ⟨rcoBroadcast | m⟩ **do**
  - **trigger** ⟨rbBroadcast | (DATA, VC, m)⟩ ← send m with VC
  - VC[self] := VC[self] + 1
  - **trigger** ⟨rcoDeliver | self, m⟩ ← VC has only increased, so RCO deliver

Fail-Silent Waiting Causal Impl. (2)

- **upon event** ⟨rbDeliver | p_j, (DATA, VC_m , m)⟩ **do**
  - **if** $p_j \neq$ self **then**
    - pending := pending ∪ (p_j, (DATA, VC_m, m)) ← put on hold
    - deliver-pending()

- **procedure** deliver-pending()
  - **while exists** x=(s_m,(DATA,VC_m,m))∈pending s.t. VC≥VC_m **do**
    - pending := pending \ (s_m, (DATA, VC_m, m)
    - **trigger** ⟨rcoDeliver | s_m, m⟩ ← Remove on hold deliver and increase local VC
    - VC[rank(s_m)] := VC[rank(s_m)] + 1

for every message whose VC_m precedes local VC

Possible Execution?

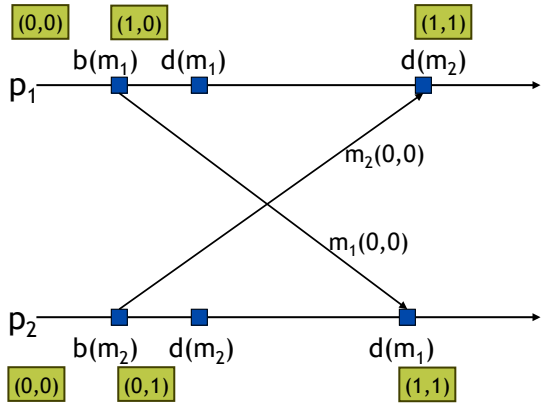

- Delivery order isn't same!
  - What is wrong? [d] Nothing, there is no causality.

Different Possible Orderings

- Some common orderings
  - (Single-source) FIFO order
  - Total order
  - Causal order

Single-Source FIFO Order

- Intuitively
  - Msgs from same node delivered in order sent
- For all messages m_1 and m_2 and all p_i and p_j,
  - if p_i broadcasts m_1 before m_2, and if p_j delivers m_1 and m_2, then p_j delivers m_1 before m_2
- Caveat
  - This formulation doesn't require delivery of both messages

Seif Haridi, based on Ali Ghodsi's slides

Total Order

- Intuitively
  - Everyone delivers everything in exact same order
- For all messages m_1 and m_2 and all p_i and p_j,
  - if both p_i and p_j deliver both messages, then they deliver them in the same order
- Caveats
  - This formulation doesn't require delivery of both messages
  - Everyone delivers in the *same order*, this might not be the *send order*!

Seif Haridi, based on Ali Ghodsi's slides

Execution Example (1)

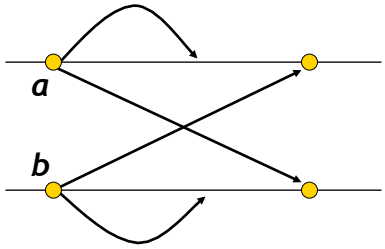

single-source FIFO? **yes**

totally ordered? **no**

causally ordered? **yes**

Execution Example (2)

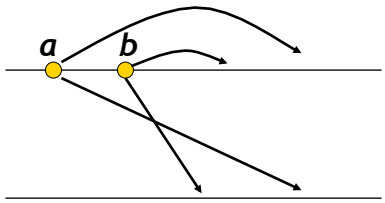

single-source FIFO? **no**

totally ordered? **yes**

causally ordered? **no**

Execution Example (3)

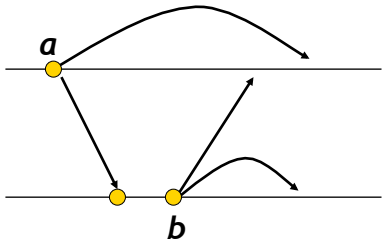

single-source FIFO? **yes**

totally ordered? **no**

causally ordered? **no**

Hierarchy of Orderings

- Stronger implies weaker ordering (→)

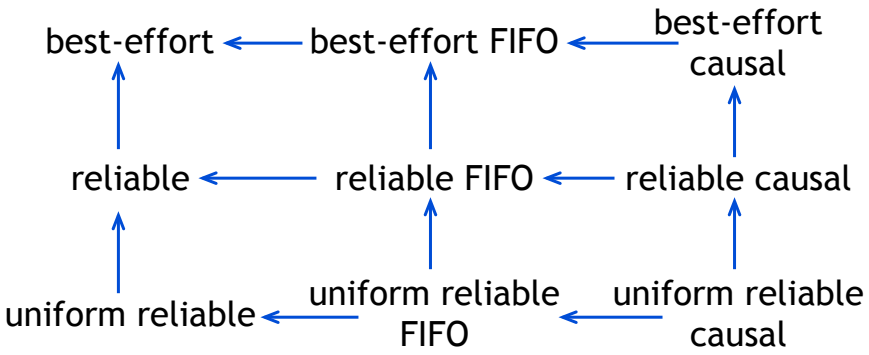

- Where does total order fit? [d]